

Summer 2022 Data Science Intern Challenge

Question 1 Analyzes

Question 1: Given some sample data, write a program to answer the following: [click here](#) to access the required data set

On Shopify, we have exactly 100 sneaker shops, and each of these shops sells only one model of shoe. We want to do some analysis of the average order value (AOV). When we look at orders data over a 30 day window, we naively calculate an AOV of \$3145.13. Given that we know these shops are selling sneakers, a relatively affordable item, something seems wrong with our analysis.

- a. Think about what could be going wrong with our calculation. Think about a better way to evaluate this data.
- b. What metric would you report for this dataset?
- c. What is its value?

Candidate Name: Dinnara Hitt

I wrote a Python program in Jupyter Notebook, using Pandas, Matplotlib, and NumPy libraries to help me analyze the data and answer question 1

```
In [1]: import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
import scipy.stats as sts

#csv file provided
csv_file_path = "2019 Winter Data Science Intern Challenge Data Set - Sheet1.csv"
```

```
In [2]: #Create a pandas dataframe with the csv file
purchase_df = pd.read_csv(csv_file_path)
purchase_df.head()
```

Out[2]:

	order_id	shop_id	user_id	order_amount	total_items	payment_method	created_at
0	1	53	746	224	2	cash	2017-03-13 12:36:56
1	2	92	925	90	1	cash	2017-03-03 17:38:52
2	3	44	861	144	1	cash	2017-03-14 4:23:56
3	4	18	935	156	1	credit_card	2017-03-26 12:43:37
4	5	18	883	156	1	credit_card	2017-03-01 4:35:11

```
In [3]: purchase_df.describe()
```

Out[3]:

	order_id	shop_id	user_id	order_amount	total_items
count	5000.000000	5000.000000	5000.000000	5000.000000	5000.000000
mean	2500.500000	50.078800	849.092400	3145.128000	8.78720
std	1443.520003	29.006118	87.798982	41282.539349	116.32032
min	1.000000	1.000000	607.000000	90.000000	1.00000
25%	1250.750000	24.000000	775.000000	163.000000	1.00000
50%	2500.500000	50.000000	849.000000	284.000000	2.00000
75%	3750.250000	75.000000	925.000000	390.000000	3.00000
max	5000.000000	100.000000	999.000000	704000.000000	2000.00000

```
In [4]: # Sort dataframe by order_amount
purchase_df = purchase_df.sort_values("order_amount", ascending=False)
purchase_df.head(100)
```

Out[4]:

	order_id	shop_id	user_id	order_amount	total_items	payment_method	created_at
2153	2154	42	607	704000	2000	credit_card	2017-03-12 4:00:00
3332	3333	42	607	704000	2000	credit_card	2017-03-24 4:00:00
520	521	42	607	704000	2000	credit_card	2017-03-02 4:00:00
1602	1603	42	607	704000	2000	credit_card	2017-03-17 4:00:00
60	61	42	607	704000	2000	credit_card	2017-03-04 4:00:00
...
3927	3928	97	979	810	5	credit_card	2017-03-11 7:37:13
2757	2758	66	772	805	5	credit_card	2017-03-14 8:43:29
3438	3439	66	842	805	5	credit_card	2017-03-22 17:58:37
742	743	12	727	804	4	cash	2017-03-14 16:38:01
1764	1765	12	789	804	4	debit	2017-03-03 3:10:50

100 rows x 7 columns

```
In [5]: # Calculating measures of central tendency for order_amount

mean_numpy = np.mean(purchase_df.order_amount)
print(f"The mean(AOV) Order Amount is {mean_numpy}")

median_numpy = np.median(purchase_df.order_amount)
```

```
print(f"The median Order Amount is {median_numpy}")

mode_scipy = sts.mode(purchase_df.order_amount)
print(f"The mode Order Amount is {mode_scipy}")
```

The mean(AOV) Order Amount is 3145.128
The median Order Amount is 284.0
The mode Order Amount is ModeResult(mode=array([153]), count=array([87]))


Conclusion:

Answer for question 1a.

- AOV is the same as the mean value.
- Because the mean(AOV) is way off compared to the median and mode
- I decided to investigate further.

Histogram for Order Amount

```
In [6]: plt.hist(purchase_df.order_amount)
plt.xlabel('Order Amount')
plt.ylabel('Counts')
plt.show()
```



Box Plot for Order Amount

```
In [7]: plt.title('Order Amount')
plt.ylabel('Amount')
plt.boxplot(purchase_df.order_amount)
plt.show()
```



Conclusion:

- The box plot shows that there is a lot of outliers for order_amount.
- Because of that, I decided to add a new column that calculates the price per item.


```
In [8]: # Add a column price_item to dataframe that is oder_amount / total_items
purchase_df["price_item"] = purchase_df['order_amount'] / purchase_df['total_items']
purchase_df.head()
```

```
Out[8]:
```

	order_id	shop_id	user_id	order_amount	total_items	payment_method	created_at	price_item	
	2153	2154	42	607	704000	2000	credit_card	2017-03-12 4:00:00	352.0
	3332	3333	42	607	704000	2000	credit_card	2017-03-24 4:00:00	352.0
	520	521	42	607	704000	2000	credit_card	2017-03-02 4:00:00	352.0
	1602	1603	42	607	704000	2000	credit_card	2017-03-17 4:00:00	352.0
	60	61	42	607	704000	2000	credit_card	2017-03-04 4:00:00	352.0

Box Plot for Price Item

```
In [9]: plt.title('Price Item')
plt.ylabel('Amount')
plt.boxplot(purchase_df.price_item)
plt.show()
```



Conclusion:

. The Box Plot for the new column price_item shows that there are some outliers. I am going to remove the outliers.

```
In [10]: # Sort dataframe by order_amount
purchase_df = purchase_df.sort_values("price_item", ascending=False)
purchase_df.head(100)
```

```
Out[10]:
```

	order_id	shop_id	user_id	order_amount	total_items	payment_method	created_at	price_item	
	2821	2822	78	814	51450	2	cash	2017-03-02 17:13:25	25725.0
	3085	3086	78	910	25725	1	cash	2017-03-26 1:59:27	25725.0
	4412	4413	78	756	51450	2	debit	2017-03-02 4:13:39	25725.0
	3167	3168	78	927	51450	2	cash	2017-03-12 12:23:08	25725.0
	490	491	78	936	51450	2	debit	2017-03-26 17:08:19	25725.0

	4882	4883	42	607	704000	2000	credit_card	2017-03-25 4:00:00	352.0
	4056	4057	42	607	704000	2000	credit_card	2017-03-28 4:00:00	352.0
	1049	1050	12	708	402	2	debit	2017-03-01 1:11:03	201.0
	3060	3061	12	990	402	2	debit	2017-03-22 23:52:30	201.0
	2258	2259	12	734	402	2	credit_card	2017-03-19 8:14:41	201.0

100 rows × 8 columns

```
In [11]: # Create a new dataframe, keeping only data with priceitem < 500
good_purchase_df = purchase_df.loc[purchase_df["price_item"] < 500]
good_purchase_df.head()
```

```
Out[11]:
```

	order_id	shop_id	user_id	order_amount	total_items	payment_method	created_at	price_item	
	308	309	42	770	352	1	credit_card	2017-03-11 18:14:39	352.0
	1929	1930	42	770	352	1	credit_card	2017-03-17 8:11:13	352.0
	979	980	42	744	352	1	debit	2017-03-12 13:09:04	352.0
	2018	2019	42	739	352	1	debit	2017-03-01 12:42:26	352.0
	4625	4626	42	809	352	1	credit_card	2017-03-11 8:21:26	352.0

```
In [12]: good_purchase_df.tail()
```

```
Out[12]:
```

	order_id	shop_id	user_id	order_amount	total_items	payment_method	created_at	price_item	
	4474	4475	92	820	180	2	credit_card	2017-03-24 12:56:45	90.0
	1927	1928	92	960	270	3	debit	2017-03-21 3:44:08	90.0
	1084	1085	92	852	180	2	cash	2017-03-24 13:44:57	90.0
	335	336	92	853	180	2	cash	2017-03-27 12:53:28	90.0
	2092	2093	92	986	90	1	debit	2017-03-04 6:44:05	90.0


```
In [13]: good_purchase_df.describe()
```

```
Out[13]:
```

	order_id	shop_id	user_id	order_amount	total_items	price_item
count	4954.000000	4954.000000	4954.000000	4954.000000	4954.000000	4954.000000
mean	2498.990916	49.819540	848.919257	2717.367784	8.851029	152.475575
std	1444.498907	29.014845	87.846007	41155.996469	116.857286	31.260218
min	1.000000	1.000000	607.000000	90.000000	1.000000	90.000000
25%	1248.250000	24.000000	775.000000	163.000000	1.000000	132.000000
50%	2494.500000	50.000000	849.000000	284.000000	2.000000	153.000000
75%	3750.750000	74.000000	925.000000	390.000000	3.000000	168.000000
max	5000.000000	100.000000	999.000000	704000.000000	2000.000000	352.000000

- You can see above that we deleted 46 rows

```
In [14]: plt.title('Order Amount')
plt.ylabel('Amount')
plt.boxplot(good_purchase_df.price_item)
plt.show()
```



In [15]: `# Calculating measures of central tendency for order_amount`

```
mean_numpy = np.mean(good_purchase_df.order_amount)
print(f"The mean(AOV) Order Amount is {mean_numpy}")

median_numpy = np.median(good_purchase_df.order_amount)
print(f"The median Order Amount is {median_numpy}")

mode_scipy = sts.mode(good_purchase_df.order_amount)
print(f"The mode Order Amount is {mode_scipy}")
```

The mean(AOV) Order Amount is 2717.3677836092047
The median Order Amount is 284.0
The mode Order Amount is ModeResult(mode=array([153]), count=array([87]))

Conclusion:

- The AOV went down to 2717.37.
- But the box plot above shows us that there are still some outliers.
- Let's remove the outliers with a Price of around \$350.00

In [16]: `#Create a another dataframe, keeping only data with priceitem < 300`
`second_purchase_df = good_purchase_df.loc[good_purchase_df["price_item"] < 300]`
`second_purchase_df.head()`

Out[16]:

	order_id	shop_id	user_id	order_amount	total_items	payment_method	created_at	price_item
1049	1050	12	708	402	2	debit	2017-03-01 1:11:03	201.0
3060	3061	12	990	402	2	debit	2017-03-22 23:52:30	201.0
2258	2259	12	734	402	2	credit_card	2017-03-19 8:14:41	201.0
1206	1207	12	721	402	2	debit	2017-03-19 13:15:17	201.0
3571	3572	12	930	402	2	cash	2017-03-22 2:02:32	201.0

In [17]: `second_purchase_df.tail()`

Out[17]:

	order_id	shop_id	user_id	order_amount	total_items	payment_method	created_at	price_item
4474	4475	92	820	180	2	credit_card	2017-03-24 12:56:45	90.0
1927	1928	92	960	270	3	debit	2017-03-21 3:44:08	90.0
1084	1085	92	852	180	2	cash	2017-03-24 13:44:57	90.0
335	336	92	853	180	2	cash	2017-03-27 12:53:28	90.0
2092	2093	92	986	90	1	debit	2017-03-04 6:44:05	90.0

In [18]: `second_purchase_df.describe()`

Out[18]:

	order_id	shop_id	user_id	order_amount	total_items	price_item
count	4903.000000	4903.000000	4903.000000	4903.000000	4903.000000	4903.000000
mean	2499.584540	49.900877	849.858862	300.155823	1.995717	150.400163
std	1444.221163	29.154367	86.887947	155.941112	0.982602	23.851202
min	1.000000	1.000000	700.000000	90.000000	1.000000	90.000000
25%	1246.500000	24.000000	776.000000	163.000000	1.000000	132.000000
50%	2499.000000	50.000000	850.000000	284.000000	2.000000	153.000000
75%	3750.500000	74.000000	925.000000	386.500000	3.000000	166.000000
max	5000.000000	100.000000	999.000000	1086.000000	8.000000	201.000000

- You can see above that we deleted 97 rows total.
- Count = 4903

In [19]: `plt.title('Price Item')`
`plt.ylabel('Amount')`
`plt.boxplot(second_purchase_df.price_item)`
`plt.show()`



In [20]:

```
# Calculating measures of central tendency for order_amount

mean_numpy = np.mean(second_purchase_df.order_amount)
print(f"The mean(AOV) Order Amount is {mean_numpy}")


median_numpy = np.median(second_purchase_df.order_amount)
print(f"The median Order Amount is {median_numpy}")

mode_scipy = sts.mode(second_purchase_df.order_amount)
print(f"The mode Order Amount is {mode_scipy}")
```

The mean(AOV) Order Amount is 300.1558229655313
 The median Order Amount is 284.0
 The mode Order Amount is ModeResult(mode=array([153]), count=array([87]))

In [21]:

```
# Histogram
plt.hist(second_purchase_df.order_amount)
plt.xlabel('Order Amount')
plt.ylabel('Counts')
plt.show()
```



Conclusion:

- After removing all outliers the mean value (what we traditionally call Average Order Value) is 300.15
- the Median(the middle value of all orders) value stayed the same throughout this process showing 284.00
- The mode (the most frequently occurring order value) value is 153.00

Final Conclusion:

- It is a good idea to use the measures of central tendency to investigate the data.
- With an AOV value so high when compared to median and mode. We could see that there was something "wrong" with the data
- After removing all the Outliers the AOV made more sense.
- If we decided not to remove the outliers the median value would be the one that made more sense to use.

Answer for question 1b. and 1c.

- Not removing the outliers - I would have used the median Order Amount of 284.0
- Removing the outliers - I think it is safe to use the mean(AOV) Order Amount of 300.15

In []: